

Cache optimizations in Garbage Collection

Bachelor Thesis

David Himmelstrup
Department of Computer Science



Research aims

- 1 Reasonable implementation complexity?
- 2 Improved performance?
- 3 Modern, fully-featured collectors?



Memory management

Top 10 programming languages (GitHub 2017):

- 1 Javascript
- 2 Python
- 3 Java
- 4 Ruby
- 5 PHP
- 6 C++
- 7 C#
- 8 Go
- 9 C
- 10 Swift




Memory management

Top 10 programming languages (GitHub 2017):


- 1 Javascript
- 2 Python
- 3 Java
- 4 Ruby
- 5 PHP
- 6 C++
- 7 C#
- 8 Go
- 9 C
- 10 Swift



Memory hierarchy




Memory hierarchy




RAM: Random-Access Memory?



Semi-space garbage collection



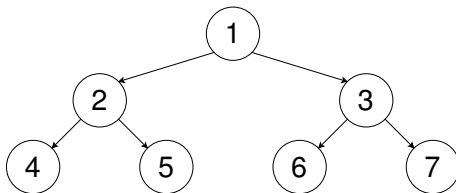
Depth-first copying




Each branch lives in continuous memory.



Breadth-first copying



Breadth-first copying



No stack but poor locality of reference.




A middle ground: tail-first copying

- 1 No stack
- 2 No reserved heap
- 3 Puts branches in *chunks* of continuous memory




A middle ground: tail-first copying



- ① No stack
- ② No reserved heap
- ③ Puts branches in *chunks* of continuous memory



New opportunity: tail compaction



Results

LLC Load Misses (millions):			
Name	Baseline	Tail-copy	Tail-compact
Allocate	3.9	+5.5%	+6.2%
BinTree	60.9	-4.5%	-4.9%
Bush	33.5	+4.1%	+2.8%
BushTail	50.9	+0.8%	+0.6%
Imbalanced	40.9	-10.3%	-11.0%
MemBench	84.1	-80.3%	-80.0%
PowerSet	130.3	-78.7%	-78.7%
Min		-80.3%	-80.0%
Max		5.5%	6.2%
Mean		-23.3%	-23.6%

- Mostly fewer LLC misses (-80% to +6%)



Results

Branch Misses (millions):			
Name	Baseline	Tail-copy	Tail-compact
Allocate	3.7	-6.1%	-13.8%
BinTree	29.0	+107.5%	+102.5%
Bush	17.4	+78.3%	+77.0%
BushTail	16.2	+44.4%	+44.2%
Imbalanced	39.4	+6.7%	+4.3%
MemBench	10.7	-7.9%	-12.8%
PowerSet	20.0	-6.7%	-10.7%
Min		-7.9%	-13.8%
Max		107.5%	102.5%
Mean		30.9%	27.3%

- Mostly fewer LLC misses (-80% to +6%)
- More complex algorithm \Rightarrow more branch misses



Results

Heap size:			
Name	Baseline	Tail-copy	Tail-compact
Allocate	1.8 gb	+0.0%	-8.3%
BinTree	1.8 gb	+0.0%	-6.2%
Bush	1.3 gb	+0.0%	-3.1%
BushTail	5.1 gb	+0.0%	-1.7%
Imbalanced	1.6 gb	+0.0%	-6.2%
MemBench	6.7 gb	+0.0%	-5.0%
PowerSet	10.4 gb	+0.0%	-5.0%
Min		0.0%	-8.3%
Max		0.0%	-1.7%
Mean		0.0%	-5.0%

- Mostly fewer LLC misses (-80% to +6%)
- More complex algorithm \Rightarrow more branch misses
- 5.0% smaller heap on average



Results

GC times:			
Name	Baseline	Tail-copy	Tail-compact
Allocate	1.5s	+14.5%	+3.1%
BinTree	4.6s	+16.2%	+10.2%
Bush	2.8s	+12.5%	+9.6%
BushTail	3.6s	+10.7%	+5.9%
Imbalanced	4.0s	-0.7%	-6.2%
MemBench	7.8s	-27.5%	-31.4%
PowerSet	12.6s	-31.1%	-34.5%
Min		-31.1%	-34.5%
Max		16.2%	10.2%
Mean		-0.8%	-6.2%

- Mostly fewer LLC misses (-80% to +6%)
- More complex algorithm \Rightarrow more branch misses
- 5.0% smaller heap on average
- Four programs are slower and three are faster.



Conclusion

- ① Reasonable implementation complexity? Yes.
- ② Improved performance? Mixed performance.
- ③ Modern, fully-featured collectors? Unlikely.

